

Lecture 1: Introduction and Overview

(Lecture Series on Foundations of Blockchains)

Focus of Lecture Series

Intended audience: has some amount of computer science training (possibly self-taught).

Focus: the science + tech of blockchains & applications built on top of them.

- not about: hype, investing, entrepreneurship, nitty-gritty engineering details
- instead: fundamental principles of block chain design + analysis

Warning:

not standardized,
in flux

"Blockchain":
stack
application layer
(stuff built on top)
layer 2
(scaling layer)
layer 1
(consensus layer)
(+ a compute layer)
layer 0
(in the Internet)

Overview of Lecture Series

Comments

Outline:

{will skip layer 0}

- ① a new computing paradigm
- ② not about digital money
- ③ principles over protocols

- first 60% of lectures about layer 1

- classical consensus protocols, Nakamoto/^{PoS()} longest-chain consensus, sybil-resistance (^{PoW}) difficulty adjustment, selfish mining, deep dives on Bitcoin, Ethereum

- next 20% on layer 2

- scaling Bitcoin via payment channels, Lightning network
- scaling Ethereum via "rollups")

- next-generation layer 1 protocols

Maybe

- last 20% on application layer

- DeFi primitives (oracles, stablecoins),

- MEV,

DAOs

DeFi apps (borrowing, lending, trading via AMMs)

Digital Signature Schemes

Consensus: (informally) keep multiple machines ("nodes") in sync, even in the face of an unreliable network, malicious attacks.

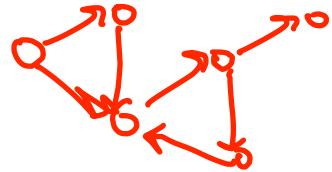
Permanent assumptions: (i) the Internet exists (ii) cryptography exists.

Definition of a DSS: (defined via 3 algorithms)

- (1) key generation algorithm (maps random seed $r \mapsto (\text{pk}, \text{sk})$ pair)
- (2) signing algorithm (maps $\text{msg} + \text{sk} \mapsto \text{msg} + \underline{\text{sig}}$)
- (3) verification algorithm (maps $\text{msg} + \text{sig} + \text{pk} \mapsto \text{"yes"/no"}$)

Assumption: ("ideal" signatures) don't know $\text{sk} \Rightarrow$ impossible to generate valid $\text{msg} + \text{sig}$.

By default: nodes sign all messages that they send.



The SMR Problem

($SMR = \frac{\text{"state machine replication"}}{\text{(1980s)}}$)

- clients submit transactions (txs) to nodes
- each node maintains a local append-only data structure (ordered sequence of txs) (*a "history"*)
- want to keep all nodes in sync, identical local histories

Definition: a **protocol** = code to be run by each node
 (local computations, receive msgs from other nodes + clients, send msgs to other nodes)

Goals: (1) Consistency: all nodes agree on history (i.e., same tx sequence)

(2) Liveness: every valid tx submitted eventually added to history

Next: the synchronous model and the Dolev-Strong protocol!