

# Lecture 7: The Tendermint Protocol

(Lecture Series on Foundations of Blockchains)

# Consensus with Partial Synchrony

Recall: Partially synchronous model. (<sup>unknown transition time  $bst$</sup>  from asynchrony to synchrony)

Goals: (for SMR) consistency (always) and liveness (eventually).

Lecture 6: if  $f \geq n/3$ , no protocol satisfies these goals (even with PKI).

Tendermint: an SMR protocol that, when  $f < n/3$ , satisfies consistency (always) + liveness (eventually). [Buchman/Kwon/Milosevic 2018]

# Tendermint: High-Level Ideas

Idea #1: iterated single-shot consensus. (output of each = ordered list of txs ("block"))  
⇒ each node  $i$  maintains its own "height"  $h_i$ ;  
(in asynchronous phase, might be different for different nodes)  
+ ignores all messages about other heights

Idea #2: for a fixed height, keep proposing + voting til agreement.

Idea #3: two stages of voting. (Necessary because different nodes may see different voting outcomes, need an intermediate "hedging" outcome.)

# Quorum Certificates (QCs)

- Preliminaries: - assume  $\text{PKI}$ , all messages signed by sender
- round = interval of  $4\Delta$  timesteps (recall shared global clock,  $\Delta$  known a priori)  
*(all nodes agree on the current round, even in the asynchronous phase)*
  - use rotating leaders (one per round) [easy due to shared clock, permissioned setting]

Definition: a quorum certificate (QC) = batch of  $\geq \frac{2}{3}n$  signed votes  
for some block  $B$  (at some height, at some round, at some stage). *Consequence: all QC for same block h + round + stage must agree on the block*

Lemma: any two QCs overlap in at least one honest node.

Proof: overlap  $\geq n - \frac{1}{3}n - \frac{1}{3}n = \frac{1}{3}n > f$ . *by assumption!*

Plan: each node  $i$  maintains a  $(B_i, QC_i)$  pair. *initially  $QC_i = \text{null}$ ,  $B_i = \text{all unexecuted txs that } i \text{ knows about}$*

- periodically updates to most recent block-QC pair it heard about  
*(also save for future use any QCs for future blocks)*

# The Tendermint Protocol

fix a height (e.g., block #9), a round  $r$  with leader  $\ell$ .

$t = 4\Delta r$ :  $\ell$  updates  $(B_e, QC_e)$  to most recent QC known, proposes to all other nodes.

$t = 4\Delta r + \Delta$ : if node  $i$  receives  $(B_e, QC_e)$  from  $\ell$ , with  $QC_e$  at least as recent as  $QC_i$ :

- broadcast first-stage vote for  $B_e$
- update  $(B_i, QC_i) := (B_e, QC_e)$
- broadcast  $(B_e, QC_e)$  to all other nodes

$t = 4\Delta r + 4\Delta$ : (just before 1st phase of round  $s+1$ )  
if know a stage-2 QC for block # $h_i$ :  
supporting a block  $B$ , commit  $B$  to local history, increment  $h_i$ ,  
(repeat as needed)

$t = 4\Delta r + 2\Delta$ : if node  $i$  receives  $\geq \frac{2}{3}n$  round- $r$  stage-1 votes for  $B$ :

- set  $QC_i := \text{this QC}$ ,  $B_i := B$
- broadcast second-stage vote for  $B_i$
- broadcast  $(B_i, QC_i)$  to all other nodes

$t = 4\Delta r + 3\Delta$ : if node  $i$  receives  $\geq \frac{2}{3}n$  round- $r$  stage-2 votes for  $B$ :

- Set  $QC_i = \text{this QC}$ ,  $B_i := B$
- Commit  $B$  to local history (1)
- broadcast  $(B_i, QC_i)$  to all other nodes
- increment  $h_i$ , re-initialize  $B_i$  &  $QC_i$

In the background (at all times): store all QCs received for blocks  $h_i+1, h_i+2, \dots$

# Tendermint: Proof of Consistency

Theorem: Tendermint Satisfies SMR consistency. (for a given block  $h$ , all honest nodes commit the same block)

Proof: Fix a height  $h$  (e.g., block #9). Let  $r = \text{1st round in which } >n/3 \text{ honest nodes cast stage-2 votes for same block } B^*$ . Call this set  $S$  prerequisite for creation of a stage-2 QC.

Intuition: nodes of  $S$  "lock in" on  $B^*$ . Because  $|S| > n/3$ , can't create any new QCs for blocks other than  $B^*$ .

Formally: (by induction) at end of round  $r$ :  
- in round  $r+1$ , no nodes of  $S$  change their mind  $\Rightarrow (i)-(iii)$  still hold at end of round

$$(i) B_i = B^* \text{ for all } i \in S$$

$$(ii) \text{QC}_i \text{ from round-}r \text{ stage-1 or later}$$

$$(iii) \text{all QCs for other blocks are from round } r \text{ or earlier}$$

$\Rightarrow$  history repeats in rounds  $r+2, r+3, \dots$  (no QCs for blocks  $B \neq B^*$ , ever again)

$\Rightarrow$  can never have stage-2 QCs that support different blocks

$\Rightarrow$  will never have different honest nodes commit to different blocks at height  $h$

QED!

# Tendermint: Proof of Liveness

Claim: Tendermint satisfies SML liveness (eventually).

Proof: Consider a transaction  $T$  known to all honest nodes.

$\Rightarrow$  fast forward to pair  $r_1, r_2$  of consecutive rounds after  $GST + \Delta$  with honest  
Lemma: at start of round  $r_1$ , every honest node working on block# $h$  or  $h+1$  [leaders  $l_1, l_2$ ]  
↑ must exist ( $f < n/3$ )

Definition: a round is **clean** if (i) post-GST (ii) honest leader (iii) all honest nodes  
working on same block# (iv) after update in 1st phase, leader's QC at least as recent as that  
at any honest node.

Lemma: clean round  $\Rightarrow$  all honest nodes commit the block proposed by leader.

Case 1: all honest nodes start round  $r_1$  working on block#  $h+1$ . (round  $r_1$  is clean, done)

Case 2: all working on block#  $h$ . Round  $r_1$  clean, commits same block. Round  $r_2$  clean, commits block  
including  $T$ .

Case 3: leader but  $n-f$  all honest nodes working on block#  $h$ .

$\Rightarrow$  round  $r_2$  clean with honest block proposal, done.

# Can We Do Better?

Main Result: if  $f < n/3$ , Tendermint satisfies consistency (always) and liveness (eventually) in the partially synchronous model.

Possible improvements?:

- can't increase # of Byzantine nodes (without compromising elsewhere)
- can't relax partially synchronous to asynchronous (" " " ")
- can't strengthen "eventually" to "always" (" " " ")
- alternative trade-offs (e.g. favor liveness over safety during an attack)
  - as in longest chain consensus (Bitcoin, Ethereum, etc.)
- same guarantees (fault-tolerance, consistency, eventual liveness) but better performance (smaller communication complexity, fewer rounds, faster recovery time post-LSST, etc.)
  - See HotStuff (Facebook Diem), Casper-ffG (Ethereum)