

Lecture 9: The Permissionless Setting and Proof-of-Work

(Lecture Series on Foundations of Blockchains)

Permissionless Consensus

Permisioned setting: all nodes known in advance. [PKI: public keys also common knowledge] PKI "trusted setup"
[aka "private blockchain" or "proof of authority (PoA) blockchain"]

Permissionless Setting:
- unknown, ever-changing node set
- clients submit transactions to nodes via broadcast
- honest nodes communicate via broadcast/gossip

Idea: extend permissioned protocols to permissionless setting?

Issue for BFT-type consensus: what's a supermajority? (don't know n)

Issue for longest-chain consensus: how to choose leader sequence? can't do round-robin
(or uniformly random selection)

Random Sampling and Sybil Attacks

Solution: (for both BFT-type & longest-chain) random sampling.

- BFT-type: choose a random "committee" of nodes from set of all nodes
- longest-chain: each round, choose a leader at random

Issue: "Sybil" attacks. (a single node can costlessly use large # of distinct IDs)

⇒ need some sybil-proof way of randomly selecting a node
↳ i.e., probability of a node's selection independent of how many distinct IDs it uses

Two Solutions:

- ① Proof-of-work (this lecture), ⇒ works best with longest-chain consensus
- ② Proof-of-stake (Lec #12). ⇒ increasingly coupled with BFT-type consensus rather than longest-chain

Consensus vs. Sybil-Resistance

Common point of confusion: consensus protocol \neq sybil-resistance mechanism.

- Consensus algorithm/protocol: among blocks proposed thus far, decide which ones get committed to honest nodes' local histories (& in what order)
- Sybil-resistance mechanism: decide which nodes get to propose/vote on blocks



[Example: "Proof-of-stake consensus" doesn't typecheck!]

Proof-of-Work

Context: need to re-deck assumptions (A1) - (A5) from lec #8

Upshot: random sampling procedure $\xrightarrow{\text{Leader selection}} i$ s.t. (under appropriate assumptions) probability of a Byzantine leader is $\leq \alpha < 1/2$.
 \Rightarrow inherit finality/linearity/chain quality guarantees from lec #8

\rightarrow proposed by [Dwork + Naor '92] for spam-fighting!

Idealll: next leader = first node to successfully solve a hard puzzle.

Assume: existence of cryptographic hash functions (e.g., SHA-256).

Random oracle (RO) assumption: CTF h behaves like a random function, with each $h(x)$ independent + uniformly distributed.



Hard puzzle: for a given threshold T , find x s.t. $h(x) \leq T$.
- e.g., if h -SHA-256, $T = 2^{128}$, looking for x s.t. $h(x)$ begins with ≥ 80 zeroes

chosen to target a fixed rate of block production
(e.g., 1/10 min in Bitcoin)

What Should Puzzle Solutions Encode?

Hard puzzle: for a given threshold Υ , find x s.t. $h(x) \leq \Upsilon$.

Rule of random oracle assumption: only way to search for a puzzle solution x is by repeated guessing (with each guess equally likely to be a solution).

- example: $h = \text{SHA-256}$, $\Upsilon = 2^{176} \Rightarrow$ expect to need 2^{80} guesses (!) before finding a puzzle solution

Question: which x 's are allowed?

Ideally: x encodes node's choices in step ②b. [Constraint: exactly one block proposal, exactly one predecessor]
I.e., for x to qualify as puzzle solution, in addition to satisfying $h(x) \leq \Upsilon$, must have the form: $x = B \parallel \text{pred} \parallel \text{pk} \parallel \text{nonce}$

"number used once"
(meaningless bits, degrees of freedom
to find x s.t. $h(x) \leq \Upsilon$)

Note: finding one puzzle solution

no help in finding any new ones. \Rightarrow Byzantine nodes can't propose > 1
block (or pred) in any given round

Proof-of-Work: Sybil-Resistance

e.g., Sun-256. Under R3, indistinguishable from a random function

Proof-of-work: given threshold τ , find x s.t. $h(x) \leq \tau$. (x has form $B||(\text{pred}||\text{pk}||\text{nonce})$)

Assume: every node searches for a solution via repeated guessing. (justified by R.3)

↳ "hashrates"

Key property: Suppose nodes $1, 2, \dots, n$ make M_1, M_2, \dots, M_n guesses per second.

\Rightarrow For each i , $\Pr[\text{node } i \text{ is leader of next round}] = \frac{M_i}{\sum_{j=1}^n M_j}$ (independent of # of pts node i uses)

Proof: by Bayes rule, $\Pr[\text{same thing}]$

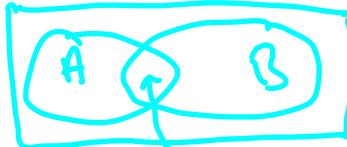
$\Pr[x \text{ proposed by node } i | x \text{ a puzzle solution}]$

$$= \frac{\Pr[x \text{ proposed by } i] \cdot \Pr[x \text{ a puzzle soln} | \text{proposed by } i]}{\Pr[x \text{ a puzzle soln}]}$$

$$= \frac{M_i}{\sum_{j=1}^n M_j} \quad \text{QED!}$$

Corollary 1: PoW is sybil proof.
[election prob. depends only on hashrates, not # of IDs]

Bayes rule



$$\Pr[A \cap B]$$

$$= \Pr[A] \Pr[B|A]$$

$$\text{also } = \Pr[B] \Pr[A|B]$$

$$\Rightarrow \Pr[A|B] = \frac{\Pr[A] \Pr[B|A]}{\Pr[B]}$$

Further Properties of Proof-of-Work

Corollary #2: if α fraction of total hashrate is Byzantine, $P[\text{first leader Byzantine}] = \alpha$.

Upshot: if < 50% of hashrate is Byzantine + assumptions $(A1) - (A5)$ hold,
PoW longest-chain consensus ("Nakamoto consensus") satisfies finality/liveness / ^{chain} quality
(with high probability)

Question: do $(A1) - (A5)$ hold?

- $(A1)$: nothing to check (trusted setup assumption)
 $(A2)$: node is the leader \iff has a valid x with $h(x) \leq T$ (which is trivial to verify)
 $(A3)$: (for fixed hashrates $\mu_1, \mu_2, \dots, \mu_n$) follow from key property/sybil-resistance
 $(A4)$: holds because round- r leader must specify block + predecessor at time of creation (1 block/round)
 $(A5)$: still need to relax (see 5th video of lecture #9) - all good provided D is small relative to average length of a round
- Note: does not require PkI 1st truly permissiveness" nor a shared global clock.

Difficulty Adjustment

Hard puzzle: given a threshold γ , find x s.t. $h(x) \leq \gamma$. $h = \text{cryptographic hash function}$
[acts like a random function]

Question: how to set γ ?

Trade-offs: (i) bigger $\gamma \Rightarrow$ faster rate of block production (Saves good for both latency + throughput)

(ii) bigger $\gamma \Rightarrow$ more frequent inadvertent forks



vs.



at least 2 of 3 will be wasted
[to quantify, need to allow $\Delta > 0 \Rightarrow$
See next video for details]
{+1; dr - all good provided γ set s.t.
avg time between solutions $\gg \Delta$ }

Answer: Set γ to achieve target rate of block production.

Case study: in Bitcoin, target rate = 1 block/10 minutes. (\Rightarrow 144 blocks/day)

\Rightarrow automatically resets γ every 2016 blocks

\Rightarrow if took $\beta \cdot (14\text{days})$ to add last 2016 blocks to the longest chain,
reset $\gamma := \gamma \cdot \beta$

ignoring details about timestamp accuracy

Work-Adjusted Longest-Chain

Issue: with varying puzzle difficulty, can't use vanilla longest-chain rule.

(adversary with modest hashrate might be able to quickly generate a long chain of blocks with very easy puzzle difficulty \Rightarrow could violate finality)

Fix: redefine chain length as total amount of hashrate that went into its construction.

- define work of block B = expected # of guesses to produce a solution at B 's difficulty level (e.g., for 256-bit hash outputs, $= \frac{2^{256}}{\gamma}$)

- work of a chain = sum of work of each block in chain

\Rightarrow honest nodes instructed to extend chain with largest amount of work

Fact: finality, liveness, etc. hold (w.h.p.) provided:

(i) at any given time, $\leq 50\%$ of hashrate is Byzantine

(ii) total hashrate only changes at bounded rate over time

(iii) security parameter k chosen appropriately

Formal analysis:

[Garay/Kiayias/Leonardos CRYPTO'17]

Crypto'17

Extensions to the Synchronous Model

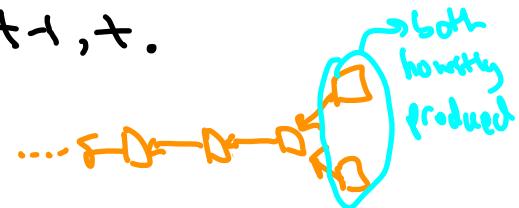
Question: what role(s) did assumption AS play in our Lecture #8 analysis?

- ① [for liveness + chain quality] implies length of longest chain \geq # honest bodies thus far every honest leader increases the length of the longest chain
- ② [common prefix property] no two honest blocks at the same height (assume fixed node hash rates for simplicity)
- ③ [for consistency] renders it trivial via instant communication

Two key parameters: $\Delta = \text{max message delay}$ $\xrightarrow{\text{equivalently, } \Pr[\text{at time } t]} = 1/L$
 $L = \text{avg amount of time between consecutive puzzle solns}$

Definition: block produced by honest node at time t is **safe** if no other puzzle solns were found in time steps $t - \Delta + 1, t - \Delta + 2, \dots, t + 1, t$.

Key point: properties ① + ② hold for safe blocks.



Consistency and Liveness (Reprise)

Lemma: expected fraction of blocks that are safe is $\geq (1-\alpha)(1 - \frac{\Delta}{L})$.

Proof: $\Pr[\text{leads to a puzzle solution}] = \Pr[\text{solution found by honest node}] \cdot \Pr[\text{in steps } \Delta+1, \dots, L]$
technically, by Union Bound $\leq \frac{1}{L}$ per fine step $\Rightarrow \leq \frac{\Delta}{L}$ per window of Δ time steps \Rightarrow prob of safe $\geq 1 - \frac{\Delta}{L}$

QED.

Note: unsafe honest blocks no worse than blocks created by Byzantine nodes.

\Rightarrow Lecture #8 analysis carries over/ provided

safe blocks > Byzantine blocks + unsafe blocks

Upshot: Nakamoto consensus ($\frac{P \cdot w}{L}$)
Satisfies consistency + liveness in general
synchronous setting with $\leq 49\%$ Byzantine
hashrate, provided Δ small relative to L .

\Rightarrow true if

$$(1-\alpha)(1 - \frac{\Delta}{L}) > \alpha + (1-\alpha) \frac{\Delta}{L}$$

$$\Leftrightarrow \alpha < \frac{1 - 2\frac{\Delta}{L}}{2 - 2\frac{\Delta}{L}}$$

PoW Can't Be Consistent Under Attack

Question 1: ever sensible to combine PoW syll-resistance + BFT-type consensus?

Question 2: can we modify longest-chain consensus to be consistent in partially synchronous setting?

Theorem: [Lewis-Pye + TR '20] No PoW blockchain satisfies both:

- (1) liveness in the synchronous setting (even with $\Delta=0$)
- (2) consistency in the partially synchronous setting.

\Rightarrow PoW + longest-chain consensus choose (1), gives up on (2)

\Rightarrow PoW + BFT-type consensus choose (2), gives up on (1)

— even assuming all nodes are honest!

→ will seem to be a dealbreaker

If idea: can't distinguish waning hash rate vs. delayed messages:

Suppose node i stops hearing new blocks from other nodes. Maybe:

- (i) other nodes turned off their machines
- (ii) other nodes' messages all delayed

Catch ->

- if node i finalizes new blocks \Rightarrow it in scenario (i), violates (2)
- if not \Rightarrow it in scenario (ii), violates (1)

QED!

Where We're Going

Lectures #10 - 13: adding a native (crypto) currency to a blockchain.

- interesting in its own right
- can incentivize nodes to run protocol (e.g., block reward)
 - worry: are nodes too incentivized (\Rightarrow selfish mining, see Lect #10)
- can charge for usage
 - need to design a transaction fee mechanism (see Lect #11)
- enables alternative approaches to sybil-resistance
 - proof-of-stake blockchains (see Lect #12)
- can help secure blockchain
 - "economic security" (\approx cost of 51% attack, see Lect #13)